

HIGH-QUALITY AI PROMPTING

Executive Desk Reference | The Master Prompt Framework

Top-performing teams treat prompts as **structured documents**, not casual questions. Use this 7-section framework to get expert-grade outputs from any AI model.

#	Section	Purpose	Example
1	Avatar	Define WHO the AI should be	"Act as a former Fortune 100 CSO with 15 years in cloud platform scaling. Avoid hype and jargon. Optimize for board-ready clarity and quantified risk."
2	Objective	State the GOAL and decision to be made	"Produce a strategic risk assessment for our proposed \$40M migration from on-prem to multi-cloud, targeting a board presentation in Q3."
3	Context	Explain WHY the problem exists	"We operate in financial services (SOC 2, PCI-DSS). Two prior migration attempts failed due to vendor lock-in and compliance gaps. Budget is fixed; timeline is 18 months."
4	Problem Space	Define SCOPE, constraints, and what has failed	"Evaluate AWS vs. Azure vs. hybrid. Prior failures: single-vendor dependency, underestimated data gravity. Compliance is non-negotiable."
5	Key Questions	Ask NON-OVERLAPPING questions that surface tradeoffs	"(a) What are the top 3 failure modes by likelihood and impact? (b) Which compliance gaps are hardest to close and why? (c) What is the realistic cost delta between hybrid and full-cloud over 5 years?"
6	Rigor	Specify HOW the AI should think	"Use first-principles reasoning. Quantify impacts where possible. Separate known facts from judgment calls. Flag unknowns explicitly. Red-team your own conclusions."
7	Output Format	Define the DELIVERABLE structure	"Produce a 2-page executive memo with: (1) decision matrix table, (2) risk heat map summary, (3) recommended path with 90-day milestones. No bullet lists longer than 5 items."

PRO TIPS: Each Key Question should require distinct reasoning and build toward a decision. Avoid yes/no framing. Surface tradeoffs, assumptions, and failure modes. Never leave output format open-ended. The more specific your format request, the more usable the output.



Prompt Framework in Action

All seven sections assembled into a single enterprise-grade prompt

AVATAR Act as a Fortune 100 CSO with 15 years in cloud platform scaling. Avoid hype. Optimize for board-ready clarity.

OBJECTIVE Produce a strategic risk assessment for our \$40M on-premises to multi-cloud migration, targeting Q3 board presentation.

CONTEXT Financial services (i.e. PCI-DSS). Prior attempts failed: vendor lock-in, compliance gaps. Budget fixed, 18-month timeline.

PROBLEM Evaluate AWS vs. Azure vs. hybrid. Non-negotiable: compliance. Prior failures: single-vendor dependency, data gravity.

QUESTIONS (a) Top 3 failure modes by likelihood/impact (b) Hardest compliance gaps to close (c) 5-year cost delta: hybrid vs. full-cloud.

RIGOR First-principles reasoning. Quantify impacts. Separate facts from assumptions. Red-team conclusions. Flag unknowns.

OUTPUT 2-page executive memo: decision matrix table, risk heat map, recommended path with 90-day milestones. Max 5-item lists.

PROMPTS & PIZZA | Master Prompt Engineering for Enterprise Teams

PLATFORM & MODEL SELECTION GUIDE

Choosing the Right AI Tool for the Job

Platform Comparison: When to Use Each AI

Capability	ChatGPT (OpenAI)	Claude (Anthropic)	Gemini (Google)
Deep Reasoning & Synthesis	Excellent (o1, o3 models)	Excellent (Opus models)	Very Good (Gemini 2.0)
Long Document Analysis	Good (128K context)	Best-in-Class (200K context)	Excellent (1M+ context)
Live Web & Search Data	Good (built-in search)	Good (web search)	Best (native Google Search)
Code Generation	Excellent	Excellent (Artifacts)	Very Good
Image Generation	Best (DALL-E 3, GPT Image)	None (no native image gen)	Very Good (Imagen 3)
Data & Spreadsheets	Very Good (Code Interpreter)	Good (file analysis)	Best (Google Sheets native)
Document Creation	Good	Very Good (Artifacts, docx)	Good (Google Docs native)
Privacy & Safety Focus	Good	Best (Constitutional AI)	Good

Model Tier Selection: Light vs. Heavy

	Light Models (Fast & Efficient)	Heavy Models (Deep & Rigorous)
Best For	Brainstorming, first drafts, summarization, translation, Q&A, simple coding tasks	Strategy, risk analysis, complex reasoning, regulatory review, multi-step synthesis, architecture decisions
Models	GPT-4o Mini, Claude Haiku, Gemini Flash	GPT o1/o3, Claude Opus, Gemini Pro/Ultra
Speed	Seconds (real-time response)	30 sec to several minutes (deep reasoning)
Cost	Low (fractions of a cent per query)	Higher (cents to dollars per query)
When to Upgrade	If output lacks depth, misses nuance, or gets facts wrong	Default here for anything going to a board, client, or regulator

QUICK DECISION RULE: If the output will influence a decision worth >\$100K, go to a client, or reach a board/regulator: use a Heavy model with a structured prompt. For everything else, start Light and upgrade if the output falls short. Always iterate: your first prompt is a draft, not a final product.

Quick-Pick Platform Guide

ChatGPT

Best for image generation, code interpreter with data viz, and deep reasoning with o-series models.

Claude

Best for long-form analysis, policy/legal docs, nuanced writing, structured document creation, and safety-critical use cases.

Gemini

Best for live web/Google data, Google Workspace integration, very long documents (1M+ tokens), and image generation.

