

PROMPTS & PIZZA

Master Prompt Engineering for Enterprise Teams

The difference between AI as a novelty and AI as a strategic asset
is a 90-minute investment in how you communicate with it.

90 Minutes • Hands-On Exercises • Enterprise Framework • Takeaway Toolkit



The AI Literacy Gap

Most teams use AI. Few use it well.



UNSTRUCTURED PROMPT

"Analyze this acquisition."

Missing: role, objective, constraints, tradeoffs, output format, rigor instructions.

Result: Generic summary. No decision support.
AI defaults to 'helpful intern' mode.



STRUCTURED PROMPT

7-part framework with role, rigor, format

Includes: expert identity, decision context, scoped constraints, analytical method, and a defined deliverable.

Result: Board-ready strategic analysis.
Repeatable. Auditable. Decision-grade.

The prompt is the interface between your strategy and the AI's execution.

Session Roadmap

90 minutes from fundamentals to enterprise-grade prompting

0–15m



The Gap

Why casual prompting fails and what replaces it

15–40m



Prompt Framework

Deep dive: 7-part Master Prompt architecture

40–55m



Prompt Surgery

Live demo: transform a weak prompt in real time

55–80m



The Lab

Hands-on: teams build enterprise-grade prompts

80–90m



Edge Plan

30-day adoption plan and shared prompt library



The Master Prompt Framework

Seven sections that turn any AI model into a domain expert

1



AVATAR

WHO the AI is

2



OBJECTIVE

GOAL and decision

3



CONTEXT

WHY the problem exists

4



PROBLEM SPACE

SCOPE and constraints

5



KEY QUESTIONS

NON-OVERLAPPING tradeoffs

6



RIGOR

HOW the AI should think

7



OUTPUT FORMAT

DELIVERABLE structure

Structure drives quality. Each section eliminates a class of AI failure modes.

01



Avatar

Define Who the AI Is

WHAT

Assign the AI a specific professional identity: role, years of experience, domain expertise, communication style, and what to avoid.

WHY IT MATTERS

Without an Avatar, the AI defaults to a generic assistant. With one, it filters every response through the lens of a domain expert.

EXAMPLE PROMPT

"Act as a former Fortune 100 CSO with 15 years in cloud platform scaling. Avoid hype and jargon. Optimize for board-ready clarity and quantified risk."



PRO TIP

Include what NOT to do. 'Avoid jargon' and 'Don't sugarcoat risks' shape tone as much as the role itself.

02 Objective

State the Goal and Decision

WHAT

Define the specific output you need and the decision it should support. Be precise about the deliverable's purpose.

WHY IT MATTERS

A clear objective prevents the AI from producing tangential analysis. It focuses reasoning on what matters to your decision-maker.

EXAMPLE PROMPT

"Produce a strategic risk assessment for our proposed \$40M migration from on-prem to multi-cloud, targeting a board presentation in Q3."



PRO TIP

Name the audience. 'For a board presentation' vs. 'for the engineering team' produces very different outputs from the same prompt.

03



Context

Explain Why the Problem Exists

WHAT

Provide the background, history, constraints, and organizational reality that shape this problem. Include what's been tried.

WHY IT MATTERS

Context prevents the AI from giving generic advice. It grounds the response in your specific situation, compliance requirements, and history.

EXAMPLE PROMPT

"We operate in financial services (SOC 2, PCI-DSS). Two prior migration attempts failed due to vendor lock-in and compliance gaps. Budget is fixed; timeline is 18 months."



PRO TIP

Include failures. Telling the AI what didn't work prevents it from recommending the same failed approach.

04 Problem Space

Define Scope and Constraints

WHAT

Bound the analysis. Specify what's in scope, what's out, what constraints are non-negotiable, and what prior approaches failed.

WHY IT MATTERS

Without boundaries, the AI will try to solve everything. Scoping produces focused, actionable analysis instead of encyclopedic overviews.

EXAMPLE PROMPT

"Evaluate AWS vs. Azure vs. hybrid. Prior failures: single-vendor dependency, underestimated data gravity. Compliance is non-negotiable."



PRO TIP

State non-negotiables explicitly. 'Compliance is non-negotiable' removes an entire branch of bad recommendations.

05 Key Questions

Surface Tradeoffs with Precision

WHAT

Ask 3–5 non-overlapping questions that each require distinct reasoning and together build toward a decision.

WHY IT MATTERS

Questions are your steering wheel. They tell the AI exactly which analytical threads to pull and ensure nothing critical is missed.

EXAMPLE PROMPT

"(a) What are the top 3 failure modes by likelihood and impact? (b) Which compliance gaps are hardest to close and why? (c) What is the realistic cost delta between hybrid and full-cloud over 5 years?"



PRO TIP

Avoid yes/no questions. Ask 'What are the tradeoffs between X and Y?' not 'Is X better than Y?'

06



Rigor

Specify How the AI Should Think

WHAT

Define the analytical method: first-principles, red-teaming, quantification requirements, confidence flagging, assumption separation.

WHY IT MATTERS

This is the section that separates amateur prompts from professional ones. It controls the depth and discipline of the AI's reasoning.

EXAMPLE PROMPT

"Use first-principles reasoning. Quantify impacts where possible. Separate known facts from judgment calls. Flag unknowns explicitly. Red-team your own conclusions."



PRO TIP

Ask the AI to 'red-team your own conclusions.' This forces it to stress-test its own logic before presenting it to you.

07



Output Format

Define the Deliverable Structure

WHAT

Specify the exact structure, length, sections, and formatting of the response. Include constraints on what NOT to produce.

WHY IT MATTERS

An open-ended format produces rambling essays. A defined format produces documents you can paste directly into a board deck or memo.

EXAMPLE PROMPT

"Produce a 2-page executive memo with: (1) decision matrix table, (2) risk heat map summary, (3) recommended path with 90-day milestones. No bullet lists longer than 5 items."



PRO TIP

Use XML tags (<analysis>, <risks>, <recommendation>) for automation pipelines. For human readers, specify sections and page limits.



Prompt Framework in Action

All seven sections assembled into a single enterprise-grade prompt

AVATAR Act as a Fortune 100 CSO with 15 years in cloud platform scaling. Avoid hype. Optimize for board-ready clarity.

OBJECTIVE Produce a strategic risk assessment for our \$40M on-premises to multi-cloud migration, targeting Q3 board presentation.

CONTEXT Financial services (i.e. PCI-DSS). Prior attempts failed: vendor lock-in, compliance gaps. Budget fixed, 18-month timeline.

PROBLEM Evaluate AWS vs. Azure vs. hybrid. Non-negotiable: compliance. Prior failures: single-vendor dependency, data gravity.

QUESTIONS (a) Top 3 failure modes by likelihood/impact (b) Hardest compliance gaps to close (c) 5-year cost delta: hybrid vs. full-cloud.

RIGOR First-principles reasoning. Quantify impacts. Separate facts from assumptions. Red-team conclusions. Flag unknowns.

OUTPUT 2-page executive memo: decision matrix table, risk heat map, recommended path with 90-day milestones. Max 5-item lists.

Live Prompt Surgery

Watch a weak prompt transform in real time

STEP 1: THE WEAK PROMPT

"Help me evaluate our cloud options."

Run it live. Show the generic output.



STEP 2: ADD THE FRAMEWORK

Apply all 7 sections live.

Let the room suggest details.

Re-run. Show the delta.

STEP 3: DEBRIEF THE DELTA

Compare the two outputs side-by-side. Ask the room:

- Which response could you send to a VP or board member?
- Which one surfaces tradeoffs you hadn't considered?
- Which one would save you 3 hours of analyst work?

Hands-On Lab

Teams of 3–4 • 25 minutes • Build a production-grade prompt



Vendor Evaluation

Evaluate 3 SaaS vendors for a \$2M enterprise platform decision. Include security, TCO, and migration risk.



M&A Due Diligence

Analyze acquisition target viability. Surface deal-breakers, synergy estimates, and integration risks.



Risk Assessment

Produce a risk assessment for a major operational change. Quantify likelihood, impact, and mitigation.



Board Strategy Memo

Draft a strategic recommendation memo for C-suite review on entering a new market or product category.

INSTRUCTIONS

Choose a scenario (or bring your own). Fill in all 7 framework sections. Run it.

Compare outputs across teams. Best prompt wins bragging rights.

Advanced Techniques

Three power moves for experienced practitioners



Few-Shot Prompting

Include 2–3 examples of ideal output directly in your prompt. This calibrates the AI's tone, depth, and format far more precisely than instructions alone.

Provide the AI with sample analyst reports showing the exact structure, terminology, and level of quantification you expect.



XML-Structured Output

Request responses in tagged sections (<analysis>, <risks>, <recommendation>). Essential for automation pipelines and ensures consistent, parseable output.

"Respond using: <thinking>reasoning</thinking> <risks>risk list</risks> <recommendation>final recommendation</recommendation>"



Iterative Refinement

Your first prompt is a draft, not a final product. After the initial response, ask: 'What information are you missing to make this analysis stronger?'

Run the prompt → Review output → Ask for missing info → Provide it → Re-run. Two iterations typically produces 3x the quality of one.

The \$100K Decision Rule

Match model capability to decision stakes

LIGHT MODELS

Fast • Low Cost • High Volume

GPT-4o Mini • Claude Haiku • Gemini Flash

- Brainstorming and ideation
- First drafts and summaries
- Translation and Q&A
- Simple code generation

HEAVY MODELS

Deep • Rigorous • High Stakes

GPT o1/o3 • Claude Opus • Gemini Pro

- Strategy and risk analysis
- Regulatory and compliance review
- Multi-step synthesis
- Architecture decisions

DECISION RULE

If the output influences a decision worth >\$100K, goes to a client, or reaches a board or regulator: use a Heavy model with a structured prompt. For everything else, start Light and upgrade if the output falls short.

Your 30-Day Adoption Plan

Structured practice builds lasting capability

WEEK 1

Foundation

Use the 7-part framework on one real task per day. Start with low-stakes requests. Save every prompt and output.

WEEK 2

Refinement

Add Rigor instructions to every prompt. Practice iteration: run, review, refine. Compare Light vs. Heavy models on the same prompt.

WEEK 3

Library

Build a team prompt library. Share your best prompts. Standardize Avatars for recurring roles (analyst, strategist, reviewer).

WEEK 4

Integration

Apply the framework to your highest-stakes work. Measure time saved. Present one AI-assisted deliverable to leadership.

PROMPTS & PIZZA

AI Amplifies Clarity. If you think clearly, it performs clearly.

The framework you learned tonight works with any model, any platform, and any level of complexity.

Your toolkit:

- Master Prompt Framework Cheat Sheet

- Executive Desk Reference (7-Part + Model Guide)

- The prompt you built tonight

